

Mesh Motion in NEMO

October 4, 2010

Abstract

NEMO has the ability to calculate mesh motion in deforming domains, so that cell quality is maximized while the computational boundaries move. In order to maintain mesh quality, vertex motion is calculated using a multi-step strategy based on optimization of surface and volume meshes. Because a single inverted cell can cause an expensive calculation to crash, this approach is designed to maximize the quality of the worst cell in the mesh. The algorithm presented here is capable of handling significant motion without connectivity changes. For the ability to handle unlimited motion, topological adaptation connectivity will be required. The following text explains the design and strategy of the mesh motion code.

1 Introduction

When CFD is required in a deforming domain, several mesh strategies are available. In static mesh approaches, the cells are truncated or body forces are used to represent walls. If, instead, the mesh moves and deforms with the domain, then the CFD code can employ standard control-volume techniques and requires only two extensions: (1) an additional flux term to represent the motion of the faces relative to the local velocity field and (2) the calculation of the change in cell volume. These relative fluxes must be calculated in a manner that is consistent with the calculation of volume in order to exactly satisfy Leibnitz's theorem ([Quan and Schmidt \(2007\)](#)). The challenge then becomes how to move vertices on surfaces and in the interior so that the mesh quality is maintained at adequate levels. The mesh moving algorithm must be extremely robust, because a single bad cell is sufficient to halt the CFD calculation.

The traditional way of calculating mesh motion is using Laplacian smoothing. In this case, we set each point to a location that is the average location of all the neighboring cells, as indicated in Eqn. 1.

$$\vec{x} = \frac{1}{N_{\text{neighbors}}} \sum_{\text{neighbors}} \vec{x}_i \quad (1)$$

On a regular mesh, Eqn. 1 is exactly the same finite difference equation that occurs if one solves Laplace’s equation for each of x , y , and z . The resulting equations can also be thought of as finding the equilibrium position of a network of springs connecting each vertex along cell edges. For a constant spring stiffness and zero equilibrium length, the resulting equation produces an unweighted averaging of vertex locations.

Laplacian smoothing is fast, simple, and easy to parallelize. Unfortunately, it is not tied directly to any measure of mesh quality. As a consequence, Laplacian smoothing frequently inverts cells, especially in tetrahedral meshes. In hexahedral meshes, even in unstructured cases, the connectivity is fairly regular. The number of cells sharing a vertex does not vary too far from the median value of eight. The angles between abutting cell edges, which act as springs in Laplacian smoothing, vary from the ideal of 90° , but not nearly to the degree found in tetrahedral meshes. In tetrahedral meshes, the number of cell edges “pulling” on a vertex varies very widely. The directions of the edge connections also vary, and will occasionally be, by chance, strongly biased to one side of a vertex location. As a consequence, Laplacian smoothing normally tangles tetrahedral meshes. The work of [Lucchini et al. \(2007\)](#), based on a finite-element smoothing process, is one of the few known examples where a form of Lagrangian smoothing has proven useful for calculations as complex as engines. They have demonstrated parallelized moving mesh computations with unstructured hexahedral meshes.

Some attempts have been made to improve the performance of Laplacian smoothing by using varying spring constants and non-zero equilibrium spring lengths ([Anderson et al. \(2005\)](#)). These methods can improve the performance of Laplacian smoothing by making spring stiffness and equilibrium length dependent on some indicator of mesh quality. This approach is not easy; if edges end up in compression instead of tension, the system of equations may be unstable or result in twisted meshes. Also, it is difficult to find a method for making spring stiffness appropriately dependent on cell quality.

The approach used here is to view the process of mesh smoothing as an optimization problem. We wish to maximize the quality of the worst cell in the mesh. This methodology is consistent with the fact that the success or

failure of a calculation is largely determined by the behavior of the worst cell. In contrast, the C++ Mesquite¹ library optimizes a smooth measure of global cell quality (Freitag et al. (2002); Brewer et al. (2003)). The trade-off is that the global optimization of a smooth measure (such as an L^2 norm of quality) is very fast, while point-wise optimization is more directly tied to the quality of the worst cells.

Ideally, the location of all the mesh vertices would be optimized simultaneously. However, finding a truly global optimum is prohibitively expensive. The point-wise approach is adopted here using the OptMS² library (Freitag (1999)). Because the surface mesh serves as a boundary condition for the interior mesh, the surface mesh is optimized first. Next, the interior mesh is optimized.

In point-wise optimization, the location of each mobile vertex is considered independently of the larger mesh. The vertex location is chosen to optimize the quality of the local sub-mesh, that consists of cells that are directly connected to the mobile, or "free" vertex. Thus, optimization is reduced to a local calculation.

A challenge for both interior and surface optimization originates from the goal of improving the quality of the worst cell. For both the interior and surface mesh smoothing, the quality function does not have a continuous first derivative, which also presents difficulties. The reason for this difficulty is that the optimization problem is of the form shown in Eqn. 2.

$$f = \max \{ \min \{ q_i(x) \} \} \quad (2)$$

Equation 2 indicates that we seek to optimize the qualities in the local sub-mesh, numbered from $i = 1$ to N_{cells} , by varying the free vertex location.. As the vertices move, the identity of the worst cell will change discontinuously. An illustration of this shift is shown in Fig. 1 for motion of a vertex, in the center of Fig. 2, in the x -direction. The figure shows a local sub-mesh of cells sharing a common vertex, which is illustrative of how the actual optimization is done. Since only the qualities of abutting cells depend on the location of a vertex, the optimization is done individually for each vertex, producing a local optimization. The position of the free vertex is chosen to maximize the minimum quality in the sub-mesh.

¹<http://www.cs.sandia.gov/optimization/knupp/Mesquite.html>

²<http://www-unix.mcs.anl.gov/~freitag/Opt-MS/>




Figure 1: The quality of several cells as the x position of the free vertex is varied. The yellow highlighted composite curve shows the minimum of all the cell qualities as a function of x , and the carret indicates the optimum.




Figure 2: A local, two-dimensional, triangular sub-mesh. The free vertex is marked with a circle.

2 General Approach

There are three phases to the mesh smoothing process. They are (1) Surface optimization-based smoothing, (2) Lagrangian smoothing of the interior, and (3) Interior optimization-based smoothing. The inclusion of Lagrangian smoothing is not necessary, but can accelerate the optimization process ([Freitag \(1999\)](#))

Because a good-quality surface mesh is necessary for maintaining the quality of the interior mesh, our process begins by optimizing the surface mesh. Where possible, the surface mesh vertices are moved along surfaces so that the surface mesh quality is optimized. This tangential movement of vertices is distinct from the Lagrangian movement of vertices on moving surfaces.

For any moving surface, Lagrangian motion of the vertices allows the surface mesh to move rigidly with the surface. The quality of the surface does not change, since, in the reference frame of the moving surface, the vertex positions are not changing. We refer to these vertices as “stick” vertices, since they stick to the moving surface.

However, in many cases, one can optimize the surface mesh by allowing vertices to slide along the surface tangent. Not all surface vertices are free to “slide.” In order to maintain the integrity of the surface, a mathematical description of the surface is required. This restriction is because curved surfaces must be projected onto flat planes for two-dimensional optimization. Once the new position of the free vertex is calculated, the new vertex location must be translated back into the corresponding location on the curved surface.

This optimization on surfaces can be considered a three-dimensional optimization that is constrained to two-dimensional curved surfaces. The mesh points are constrained to slide along surfaces, but unless the surface happens to be a plane, the tangent of the surface will vary with the location of the vertex. Thus surface mesh optimization is more complex than in the interior optimization, since the directions in which points can move is a function of where the points are currently located.

If significant discretization error occurred during the calculation of the new vertex position, then the location of the boundary vertices could drift from the true surface over time. If the surface is recognized as a plane or cylinder, the vertex location can be reconciled with the canonical surface ([Freitag et al. \(2002\)](#)). So unless a precise mathematical description of the surface is available, then the surface points are constrained to stick to the surface. If the surface can be described as a cylinder or plane, then the points can slide relative to the surface and yet remain precisely on the




Figure 3: The mesh points, on a cylinder surface, are projected to the local tangent plane. The plane is defined by the tangent at the location of the free vertex, denoted by the X.

surface. An example of optimizing “sliding” vertices is shown in Fig. 3.

In the example shown in Fig. 3, the free vertex and surrounding vertices will be projected to the tangent plane. The free vertex’s location will be optimized, which moves it slightly off the cylindrical surface. The vertex will then be moved in the direction normal to the cylinder, placing it back on the cylinder surface. This process lets vertices slide without degrading the representation of the surface.

In order to reduce the required effort and likelihood of mistakes in performing moving mesh calculations, we have implemented a very convenient “automatic surface recognition” algorithm. This algorithm, with no user inputs other than the mesh, recognizes canonical surface shapes, such as a plane or cylindrical surface. The surface need only conform to any portion of an infinite plane or cylinder (e.g. the algorithm recognizes a curved surface that might be only twenty degrees of a cylindrical surface). The algorithm uses a least-squares fitting process to test every surface.

For a plane, the least-squares fit requires the solution of a linear system of four equations. The procedure for fitting more general surfaces requires an iterative approach. In the code, treatment of a more general surface is demonstrated by recognizing cylindrical surfaces. A general cylinder equation is represented by an infinite cylindrical surface with rotation and translation. The Levenberg-Marquardt algorithm is used to find the cylinder orientation, location, and radius that best fits the surface points ([Vandevender and Haskell \(1982\)](#)). This fitting is performed only once; if a surface moves during the calculation, the constants representing the cylinder or plane position are updated analytically.

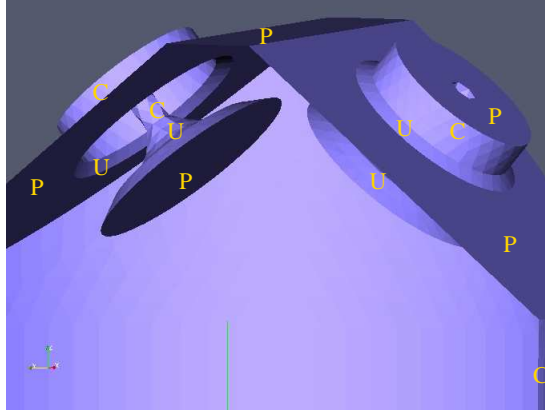


Figure 4: A cut-away view of a cylinder and four-valve engine showing automatically recognized surfaces.. Cylindrical surfaces are denoted with "C", planar with "P", and unrecognized with "U."

Currently, only recognition of planes and cylinders have been implemented, however extensions for spheres or cone frustra are possible. If the surface is neither a plane or a cylinder, then the algorithm tags the surface as "unrecognized" and moves on. An example is shown in Fig. 4. Only recognized surfaces may have sliding vertex optimization, since the surface normal of an unrecognized surface cannot be known exactly.

After the optimization of the surface mesh, the interior mesh is optimized. This can be done with direct application of analogous methods to those used in the two-dimesional phase. However, it may be expedient to use Lagrangian smoothing to provide a provisional mesh solution. The results of the Lagrangian simulation are adequate for the great majority of the cells. A small percentage of the cells will be tangled or very low quality. The optimization phase of the interior phase follows. The interior mesh is untangled as it is optimized using several iterative sweeps through the interior. The final result will be a high-quality mesh, so long as the limitations of the algorithm are respected.

The ability to untangle meshes is fairly uncommon and is quite helpful. Tangled meshes have negative volumes and present difficult mathematical problems for optimization algorithms. The current approach for untangling is based on a linear programming method to find a location of the free vertex where all cell volumes are positive (Freitag (1999)). If the local sub-mesh is sufficiently tangled, a solution may not exist and the untangling method will fail. However, by using repeated sweeps through the domain, a region of tangled cells will gradually be untangled, as the untangling algorithm finds valid vertex locations for the less tangled cells.

The untangling capabilities of the code are quite important. This means that invalid initial meshes can be repaired and, if inverted cells occur as a result of the Lagrangian vertex solution, they can be repaired. The outcome of the untangling process is usually a low-quality sub-mesh, but using several optimization passes will usually bring the mesh up to an acceptable quality.

3 Code Structure

The NEMO code, in which we have demonstrated the moving mesh capability, is an inherently parallel software library built using object-oriented Fortran 95 (Toninel (2006)). NEMO is designed as a set of foundation classes upon which specific applications can be built, similar to the approach of OpenFOAM³ (Weller et al. (1998)). Much of the lower level mesh smoothing for surfaces and interior points is based on a modified version of the OptMS library developed by Freitag (Freitag (1999)). The resulting code is fully parallelized, after an initial domain decomposition, using an “owner computes” paradigm.

Initially, the mesh is read in by all processors, and any recognizable surfaces are noted. Data structures, unique to the PSBLAS library, are initialized at this point. These data structures facilitate parallel communication about mesh entities and allow solutions of equations at cell centers, verticies, or face centers. Decomposition is handled by the ParMetis⁴ library (Karypis et al. (2003)), after which, each processor only holds the local mesh in memory, with halo layers of cells, faces, and vertices. This halo layer, which can be of arbitray depth, provides a copy of the recent state of the mesh that lies on neighboring processors and abuts the processor boundary. The halo management is based on the PSBLAS⁵ library (Filippone and Colajanni (2000); Buttari (2006)), which can keep data synchronized and exchange information with a higher-level interface than MPI provides.

As time advances, first the surfaces are moved with the associated surface mesh. Next, the sliding surface meshes are optimized using the two-dimensional capabilities of the OptMS library. Repeated sweeps are made over the surface as the optimization is done point-wise. Laplacian smoothing then provides an initial guess of the interior vertex positions. Several optimization passes then completes the mesh optimization. The steps, as described in detail below, are listed here:

1. Move surfaces and associated points

³<http://www.opencfd.co.uk/openfoam/>

⁴<http://www-unix.mcs.anl.gov/~freitag/Opt-MS/>

⁵<http://www.ce.uniroma2.it/psblas/>




Figure 5: The projection of a point to a tangent plane, as seen from the side. The red line shows the trigonometric projection and the blue arc shows the length preserving projection.

2. Untangle and optimize surface meshes
3. Lagrangian interior smoothing
4. Interior untangling and optimization

Each of these steps is described in detail below.

3.1 Moving Surfaces and Associated Points

When a surface, such as piston face or valve, moves, the points that lie on that surface should move as well. In general, the points are moved in a Lagrangian fashion. If a surface is of a recognizable type, such as a plane or cylinder, the mathematical representation of the plane and cylinder are moved as well.

3.2 Untangle and Optimize Surface Meshes

The sliding mesh optimization first projects the vertices to a tangent plane. Unlike a typical trigonometric projection, a length-preserving projection is used, as shown in Fig. 5. The length-preserving projection better represents the edges in the plane. The x and y positions of each vertex in the local surface sub-mesh are calculated in a local two-dimensional coordinate system that has the origin located at the free vertex.

Because the OptMS library only supports optimization of simplices (triangles and tetrahedra), the algorithm checks first to see if the free vertex is

shared only by triangles. For any other kinds of cells, the vertex location, in the tangential plane, is set to the average location of the neighboring vertices. Because optimization must be preceded by untangling if any cells in the sub-mesh are inverted, the algorithm also checks that all the cells in the surface sub-mesh are valid. If any local cells are tangled, then untangling is applied, otherwise the optimization routine is called. The OptMS library applies optimization to the vertex using a scheme analogous to steepest descents, but generalized for the non-smooth problem represented in Fig. 1.

3.3 Lagrangian Interior Smoothing

Once the surface mesh is smoothed, the interior vertex positions are calculated using Laplacian smoothing. This calculation is performed separately for the x , y , and z coordinates, which are each independent equations. The matrix is exactly the same for all three equations, so the matrix and preconditioner data are re-used. The parallel conjugate gradient solver in PS-BLAS, with ILU preconditioning, typically converges in about five to ten iterations. The resulting mesh often contains a small number of unacceptable cells, but the computational cost is very low, and it provides a good estimate of the final mesh. Unlike the local point-wise optimization, the Lagrangian smoothing communicates information globally.

The boundary conditions used in the Laplacian smoothing are all Dirichlet. It is assumed that all surface points will remain where the surface smoothing put them. If the interior mesh vertices lie at a boundary between two different kinds of cells, such as tetrahedra and hexahedra, then the vertices are also fixed during this stage. The reason for fixing these interior vertices is that Lagrangian smoothing is overly biased by the number and orientation of edges. At the transition between hexahedral and tetrahedral cells, the greater number of edges present in the tetrahedral region will pull the vertices towards the tetrahedral region, which may not be desirable. As a consequence, the boundaries between two kinds of cells are not permitted to move in the current implementation.

3.4 Interior untangling and optimization

The final phase of optimization is the point-wise sweeps over the mesh interior. Depending on the kind of cells and the validity of the local sub-mesh, several different treatments of the vertex are applied. Consistent with the Laplacian smoothing phase, when vertices are connected to a mixture of cell types, the vertex is left in place. For vertices used by non-tetrahedral cells,

simple averaging of locations is performed, equivalent to Laplacian smoothing.

For vertices that are used only by tetrahedra, true optimization is applied. Because optimization must be preceded by untangling when any cells in the sub-mesh are inverted, the algorithm first checks that all the cells in the sub-mesh are valid. If any local cells are tangled, then untangling is applied. For valid local sub-meshes, containing only tetrahedra, three-dimensional optimization using the OptMS library is performed. Between four to ten sweeps through the mesh interior are sufficient to achieve an optimally smoothed mesh.

References

- ANDERSON, A., ZHENG, X., AND CRISTINI, V. 2005. Adaptive unstructured volume remeshing - i: The method. *Journal of Computational Physics*.
- BREWER, M., DIACHIN, L., KNUPP, P., LEURENT, T., AND MELANDER, D. 2003. The Mesquite Mesh Quality Improvement Toolkit. In *Proceedings of the 12th International Meshing Roundtable*. Santa Fe NM, 239–250.
- BUTTARI, A. 2006. Software Tools for Sparse Linear Algebra Computations. Ph.D. thesis, University of Rom “Tor Vergata”.
- FILIPPONE, S. AND COLAJANNI, M. 2000. PSBLAS: a Library for Parallel Linear Algebra Computations on Sparse Matrices. *ACM Transactions on Mathematical Software* 26, 4, 527–550.
- FREITAG, L. 1999. Users Manual for Opt-MS: Local Methods for Simplicial Mesh Smoothing and Untangling. ANL Technical Report ANL/MCS-TM-239, Argonne National Laboratory, Argonne, Illinois. March.
- FREITAG, L., KNUPP, P., LEURENT, T., AND MELANDER, D. 2002. MESQUITE Design: Issues in the Development of a Mesh Quality Improvement Toolkit. In *Proceedings of the 8th Intl. Conference on Numerical Grid Generation in Computational Field Simulations*. Honolulu, 159–168.
- KARYPIS, G., SCHLOEGEL, K., AND KUMAR, V. 2003. *ParMETIS – Parallel Graph Partitioning and Sparse Matrix Ordering Library – Version 3.1*.

- LUCCHINI, T., D'ERRICO, G., JASAK, H., AND TUKOVIC, Z. 2007. Automatic Mesh Motion with Topological Changes for Engine Simulation. *SAE Technical Papers 2007-01-0170*.
- QUAN, S. AND SCHMIDT, D. 2007. A moving mesh interface tracking method for 3d incompressible two-phase flows. *Journal of Computational Physics*.
- TONINEL, S. 2006. Development of a New Parallel Code for Computational Continuum Mechanics Using Object-Oriented Techniques. Ph.D. thesis, University of Bologna.
- VANDEVENDER, W. H. AND HASKELL, K. H. 1982. The SLATEC mathematical subroutine library. *SIGNUM Newsl.* 17, 3, 16–21.
- WELLER, H. G., G. TABOR, JASAK, H., AND FUREBY, C. 1998. A Tensorial Approach to Computational Continuum Mechanics Using Object-Oriented Techniques. *Computers in Physics* 12, 6.